

# Resource Synchronization Protocols

Embedded OS Implementation

Prof. Ya-Shu Chen  
National Taiwan University  
of Science and Technology

# Motivation

- Access control to resources are applied to secure the validity
  - To avoid race conditions
- A resource can be shared by a number of tasks
  - Which task should be granted for access and which should not be granted?
    - The waiting time must be controllable and predictable
  - Resource sync protocol v.s. resource scheduling

# Resources

- Let us consider **passive** resources only
  - A task is still busy on CPU when it is using a passive resource
  - They do not suspend themselves on CPU
  - E.g., memory protected by semaphores
- Active resources are not considered here
  - A task may becomes idle on CPU and switch to an active resource
  - E.g., I/O devices
  - Active resources, of course, can be modeled as passive resources with loss of parallelism

# Resources

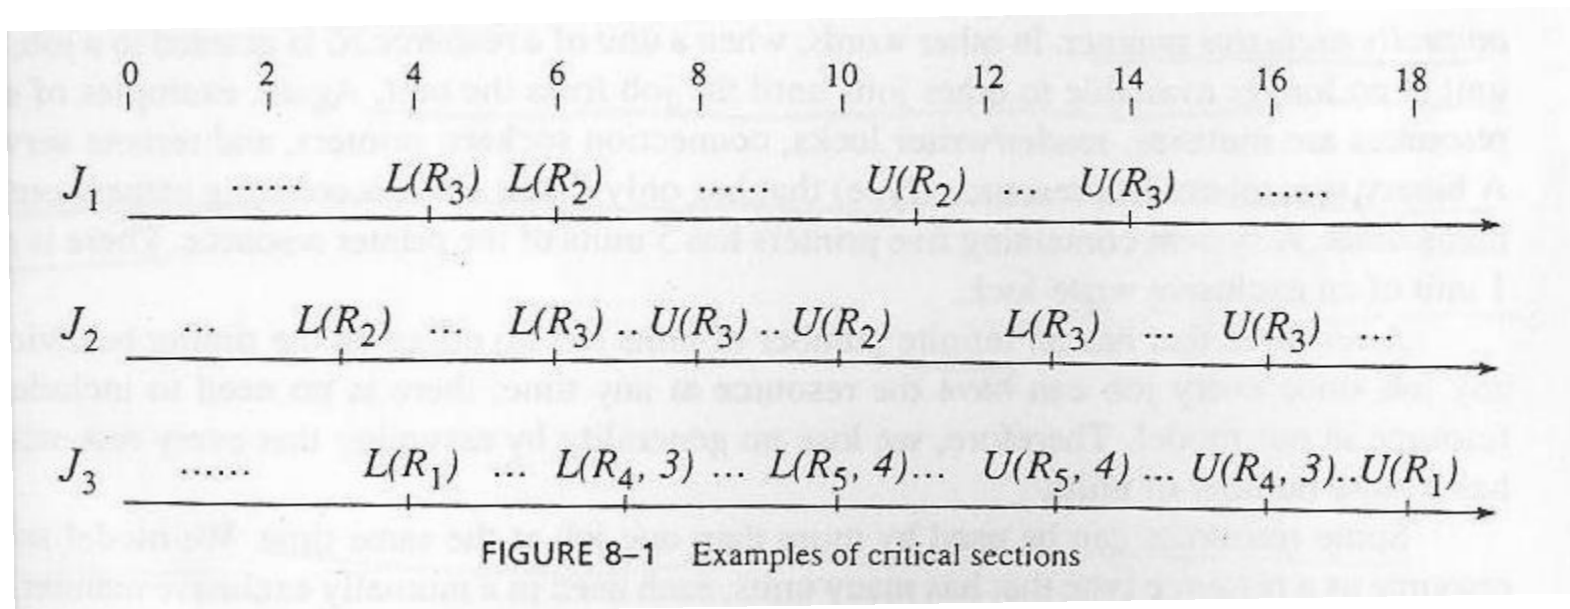
- Consider there are  $p$  types of serially reusable resources  $R_1, R_2, \dots, R_n$
- For some  $R_i$ , there are  $v_i$  indistinguishable units of resource
  - A read lock may permit many readers
  - A write lock allows only one writer
    - And also prohibit any other readers

# Resources


- When a job attempts to use a resource, we say that it tries to **lock** it
  - The lock may not be granted immediately
  - The task becomes **waiting** until the lock is granted
    - A lock is not granted because of 1) there are no such kind of resources available, or 2) to prevent some undesirable effects
- When a job no longer needs a resource, it **unlocks** it
  - An unlock is immediately granted, naturally

# Resources

- Let the duration between the lock and unlock to a resource be a **critical section**
- Let critical sections be **properly nested**





# Resources



- Resource usages may postpone the completion of tasks
  - If  $J_1$  and  $J_2$  do not require resources, they complete by 11 and 14, respectively

# Resources

- A job is said to be directly blocked by another job if it loses in the contending for some resource
- Use wait-for graph to describe run-time resource usage
- Use a bipartite graph to describe resource requirements






# Undesirable Effects

- Priority inversion
  - A high-priority job can not run because of the interference from a low-priority job
    - If the high-priority job tries to lock a resource that is currently locked by another low-priority job
  - In this case, the high-priority job is said to be **blocked**
    - Note: a low-priority job never be “blocked” by high-priority jobs (why?)

# Undesirable Effects



- Priority inversion
  - Priority inversion could damage timeliness!
    - A job may be blocked by another job, even if the two jobs do not have any direct resource contention!!



If there are many such jobs, how long  $J_1$  is blocked becomes unpredictable!

# Undesirable Effects

- Timing anomalies
  - Ideally, if the processing of a job (or a part of it) is accelerated, all jobs in the system enjoy shorter (or equal) response time
  - However, it is not true if jobs contend for resources



Now the duration of this CS becomes shorter...

# Undesirable Effects

- Deadlocks
  - The most undesirable effect to real-time systems
  - If jobs are in circular waiting, then deadlocks occur
  - Must find a way to impose a partial order on run-time resource access

# Undesirable Effects

- Priority inversions and timing anomalies are spontaneous if there are accesses to resources
  - Can durations for priority inversions be **bounded**?
  - Can timing anomalies be **controllable**?
- Deadlocks must **be avoided**
  - A system suffers from deadlock is inherently incorrect!


# Resource Synchronization Protocols

- A set of rules to
  - grant a lock to a resource
  - schedule jobs that use resources

# Non-Preemptible Critical Section


- [Mok] **NPCS**: If a job is using a resource, it won't be preempted by any other jobs
  - Even if there are no resource conflicts
  - Effectively the job runs at the highest priority
- **Deadlocks** never occur because any job holding resources can not be preempted
- **Uncontrollable priority inversions** never occur (why?)

# Non-Preemptible Critical Section



← Ungoverned

NPCS→



\*\*Not the same system :P



# Non-Preemptible Critical Section

- Pros:
  - Easy to implement
  - No need to know resource usage *a priori*
- Cons:
  - Low concurrency
  - Poor responsiveness

# Non-Preemptible Critical Section

- Let tasks in  $\{T_1, T_2, \dots, T_n\}$  be sorted in the rate-monotonic order and scheduled by **RMS**
- The longest blocking time suffered by  $T_i$  because of resource contention is

$$\max_{i+1 \leq k \leq n} (c_k)$$

- $C_k$  stands for the longest critical section of task  $T_k$

Why at most once?

# Non-Preemptible Critical Section

- If tasks are scheduled by **EDF**, the longest blocking time suffered by  $T_i$  because of resource contention is

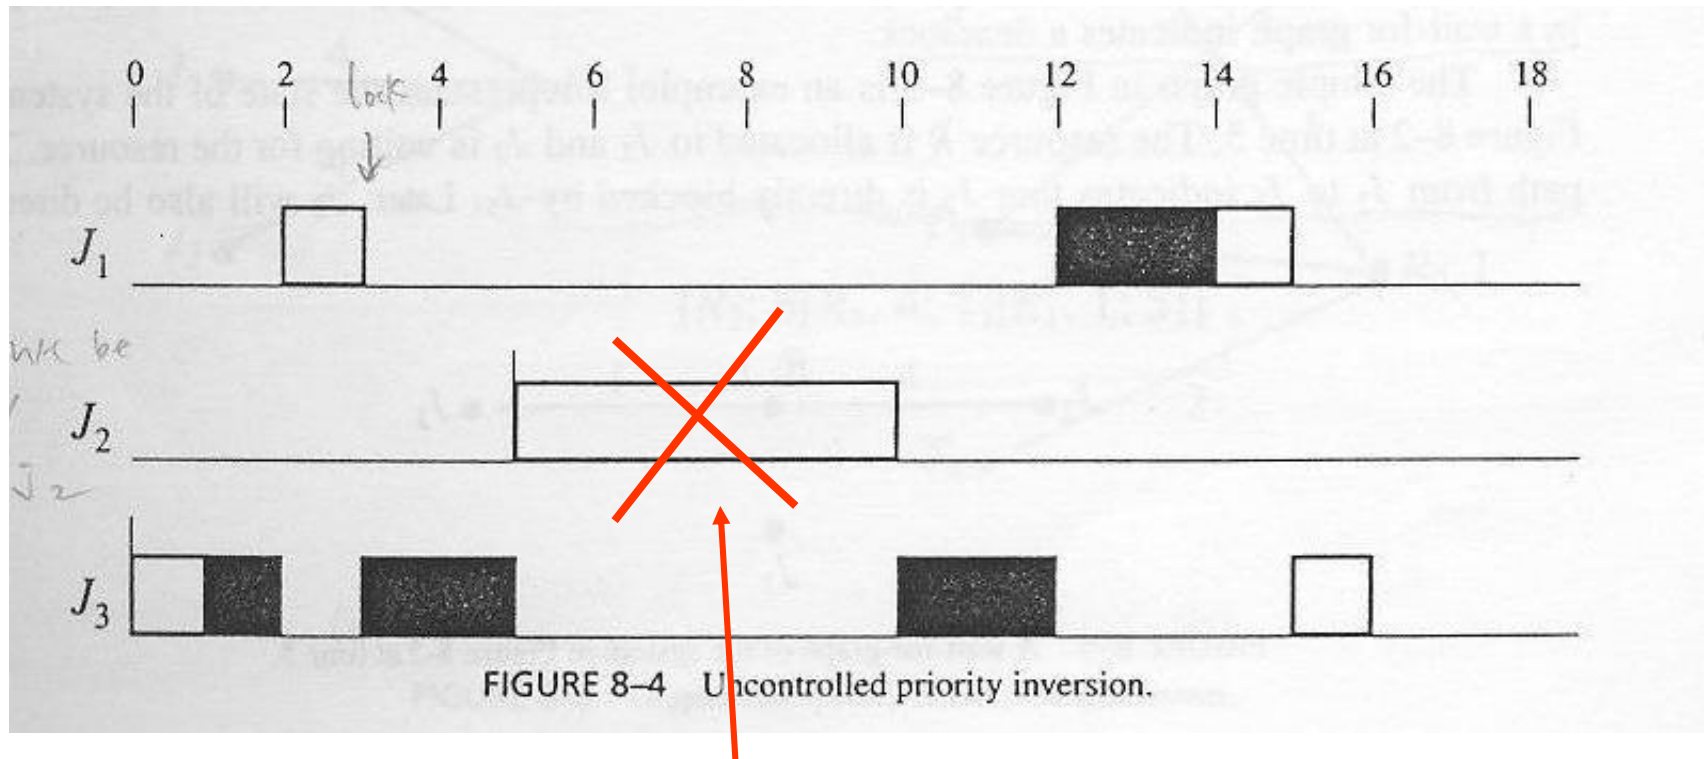
$$\max_{i+1 \leq k \leq n} (c_k)$$

- Still the same (why ?)

# Priority-Inheritance Protocol



- NPCS is **too** restrictive, it prevents high-priority jobs from executing even though they have no resource conflicts
  - CPP is restrictive too, jobs may be blocked before they try to lock resources
- PIP relieve such restriction under some certain conditions
  - uncontrollable blocking must not occur

# Priority-Inheritance Protocol



This job should not run because a higher-priority job is already waiting!!

Can J3 inherit J1's priority? Since J3 blocks J1, logically in the system there is a high-priority job to be executed



# Priority-Inheritance Protocol


## *Rules of the Basic Priority-Inheritance Protocol*

1. *Scheduling Rule*: Ready jobs are scheduled on the processor preemptively in a priority-driven manner according to their current priorities. At its release time  $t$ , the current priority  $\pi(t)$  of every job  $J$  is equal to its assigned priority. The job remains at this priority except under the condition stated in rule 3.
2. *Allocation Rule*: When a job  $J$  requests a resource  $R$  at time  $t$ ,
  - (a) if  $R$  is free,  $R$  is allocated to  $J$  until  $J$  releases the resource, and
  - (b) if  $R$  is not free, the request is denied and  $J$  is blocked.
3. *Priority-Inheritance Rule*: When the requesting job  $J$  becomes blocked, the job  $J_l$  which blocks  $J$  inherits the current priority  $\pi(t)$  of  $J$ . The job  $J_l$  executes at its inherited priority  $\pi(t)$  until it releases  $R$ ; at that time, the priority of  $J_l$  returns to its priority  $\pi_l(t')$  at the time  $t'$  when it acquires the resource  $R$ .

Note: a job may inherits priority from one or more jobs (nested inheritance)

Job	$r_i$	$e_i$	$\pi_i$	Critical Sections
$J_1$	7	3	1	[Shaded; 1]
$J_2$	5	3	2	[Black; 1]
$J_3$	4	2	3	
$J_4$	2	6	4	[Shaded; 4 [Black; 1.5]]
$J_5$	0	6	5	[Black; 4]

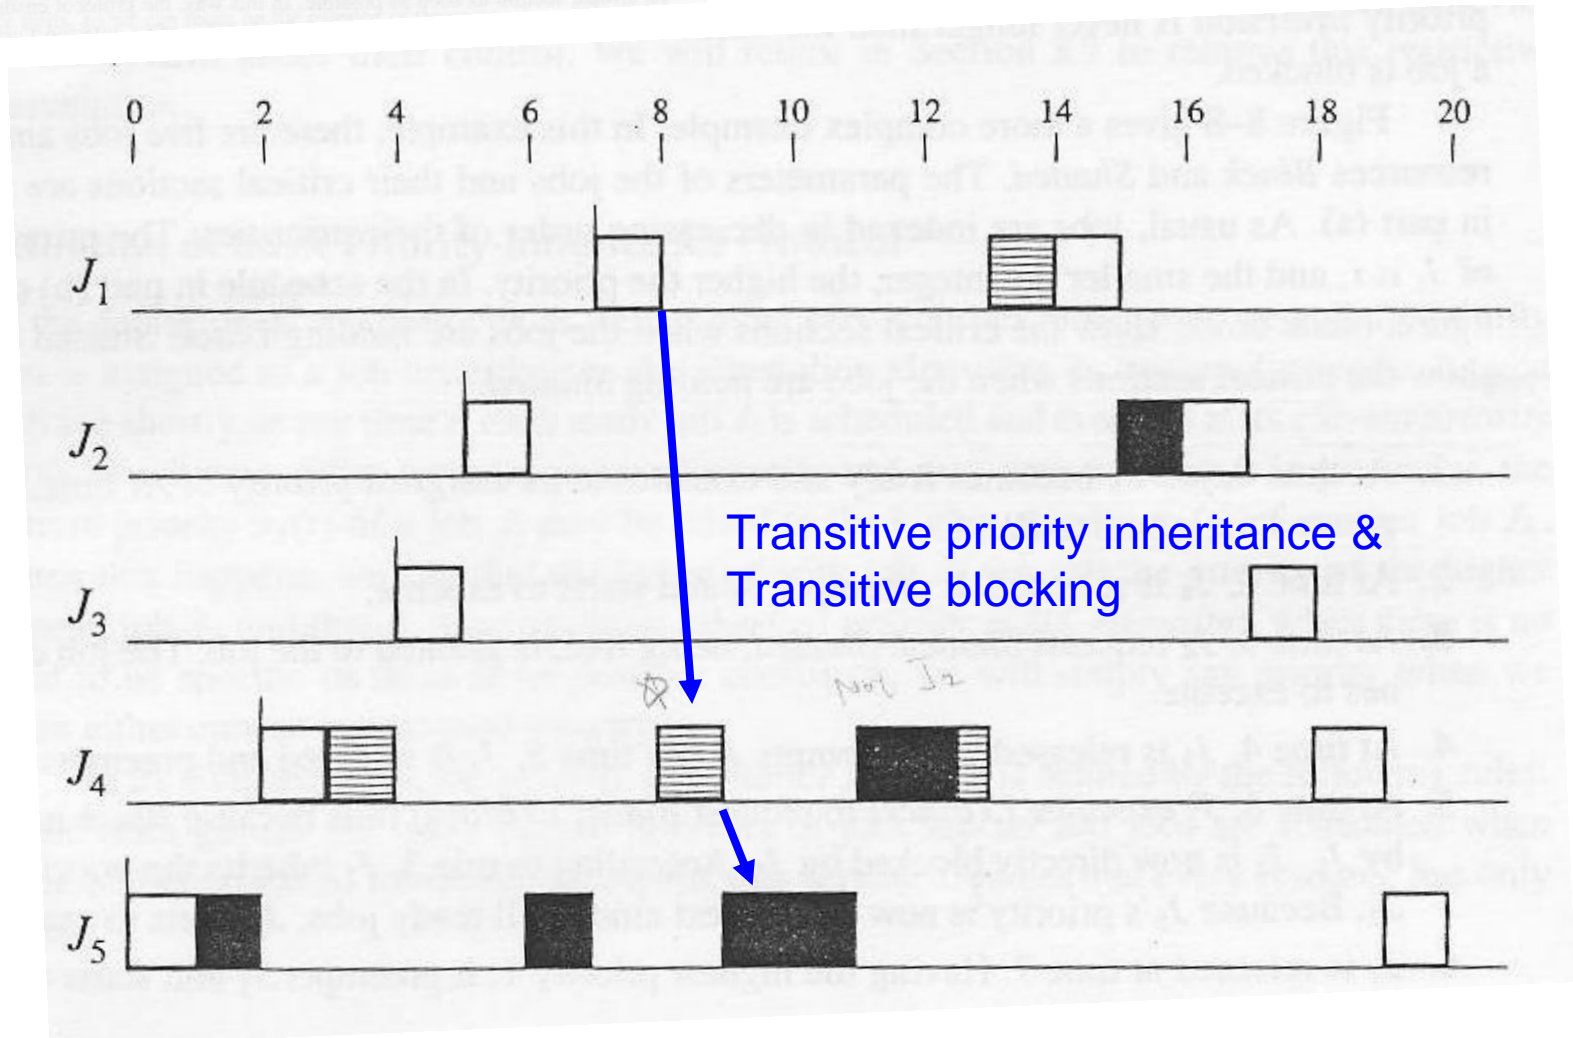
(a)






Job	$r_i$	$e_i$	$\pi_i$	Critical Sections
$J_1$	7	3	1	[Shaded; 1]
$J_2$	5	3	2	[Black; 1]
$J_3$	4	2	3	
$J_4$	2	6	4	[Shaded; 4 [Black; 1.5]]
$J_5$	0	6	5	[Black; 4]

(a)




# Priority-Inheritance Protocol

- Is chain blocking prevented?




# Priority-Inheritance Protocol



$J_1$  can be blocked up to  $\min(v, k)$  times, each of the duration of the outmost CS.  $v$  and  $k$  stand for different resources  $J_1$  requires and the number of low-priority tasks, respectively


# Priority-Inheritance Protocol

- Deadlocks




# Priority-Inheritance Protocol

- In PIP, blocking time might be lengthy (still bounded though...)




# Priority-Inheritance Protocol



NPCS

Actually more high-priority jobs can run, provided there is no resource conflicts



PIP

# Priority-Inheritance Protocol

- Pros:
  - Simple
  - High concurrency (high responsiveness)
- Cons:
  - Suffering from deadlocks
  - Blocking time is not well managed
    - Chain blocking
    - Transitive blocking

# Priority-Ceiling Protocol


- PIP avoids uncontrolled priority inversions, but jobs suffer still from
  - Deadlocks
  - Lengthy blocking time (chain blocking, transitive blocking)
- PCP aims at resolving the shortcomings of PIP



# Priority-Ceiling Protocol

- Assumptions
  - All tasks have unique and fixed priorities
  - Resource usages of tasks are known *a priori*
- Terms
  - $\Pi(R)$ : priority ceiling of resource R
    - The highest priority of all tasks that require R
  - $\Pi^{\wedge}(t)$ : current system ceiling at time t (*caution!*)
    - The highest priority ceiling of all resources that are in use

# Priority-Ceiling Protocol



# Priority-Ceiling Protocol

## Rules of Basic Priority-Ceiling Protocol

### 1. Scheduling Rule:

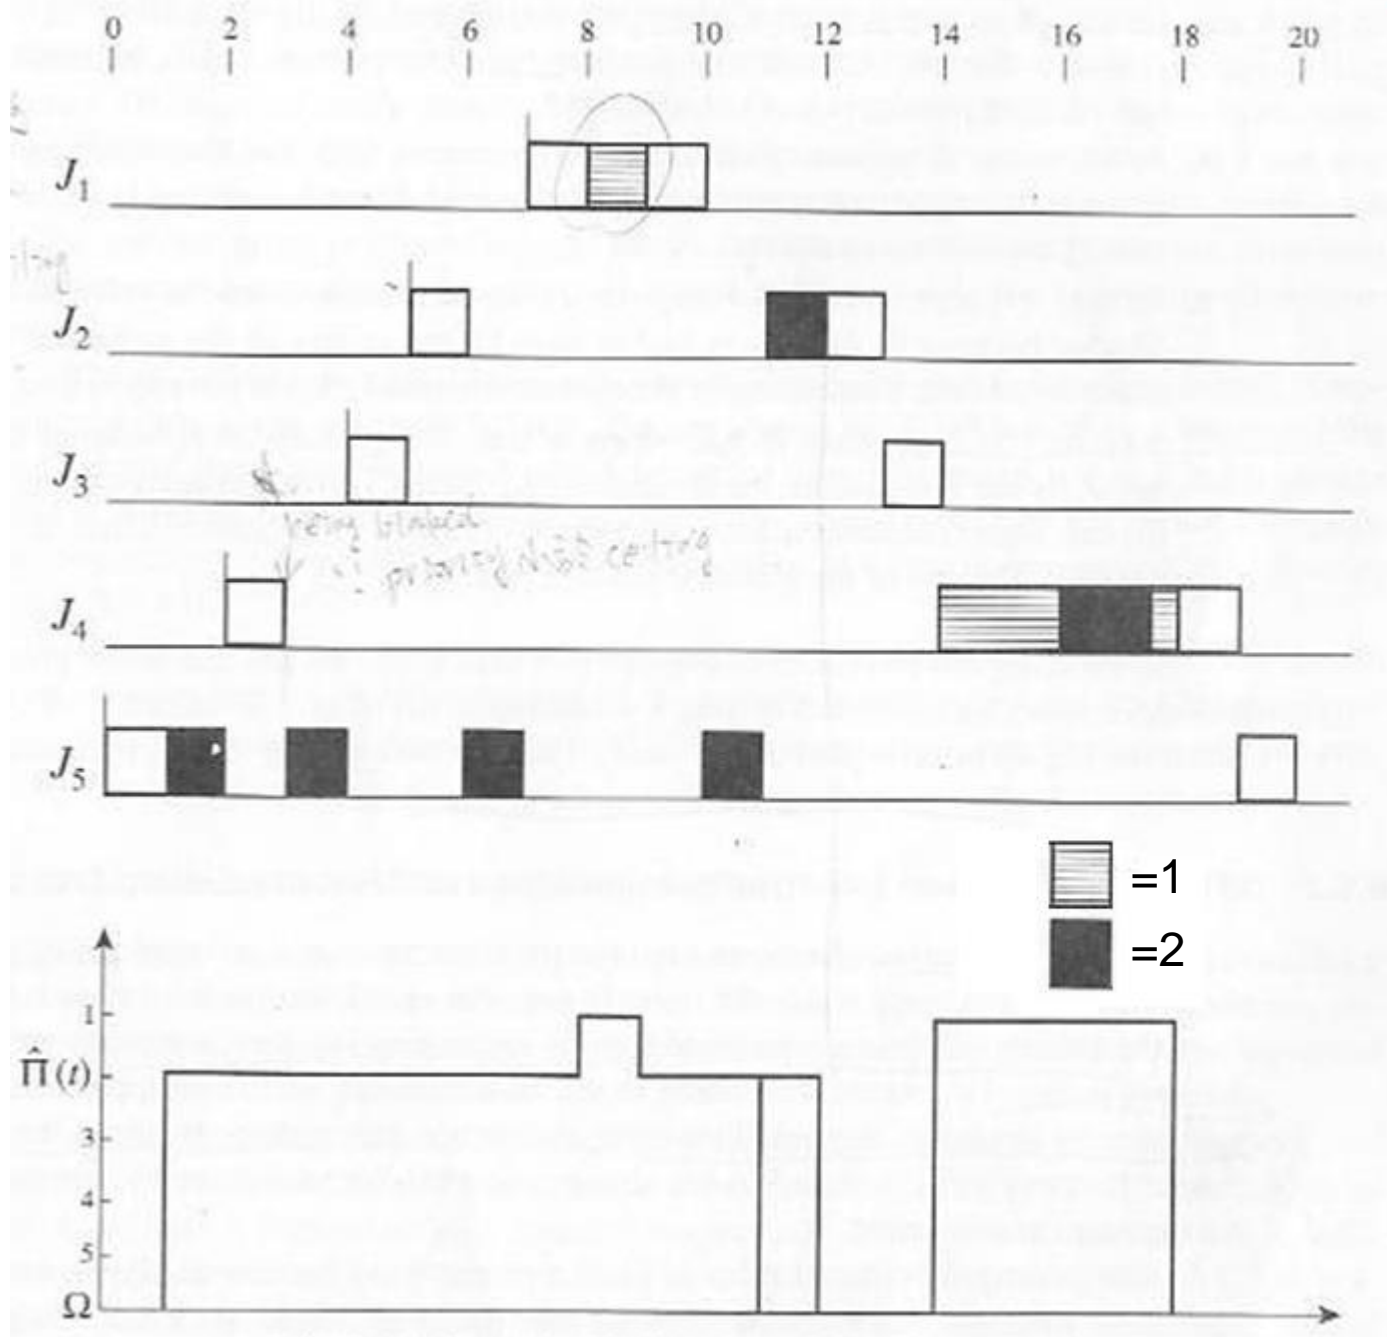
- (a) At its release time  $t$ , the current priority  $\pi(t)$  of every job  $J$  is equal to its assigned priority. The job remains at this priority except under the condition stated in rule 3.
- (b) Every ready job  $J$  is scheduled preemptively and in a priority-driven manner at its current priority  $\pi(t)$ .

### 2. Allocation Rule: Whenever a job $J$ requests a resource $R$ at time $t$ , one of the following two conditions occurs:

- (a)  $R$  is held by another job.  $J$ 's request fails and  $J$  becomes blocked.
- (b)  $R$  is free.
  - (i) If  $J$ 's priority  $\pi(t)$  is higher than the current priority ceiling  $\hat{\Pi}(t)$ ,  $R$  is allocated to  $J$ .

- (ii) If  $J$ 's priority  $\pi(t)$  is not higher than the ceiling  $\hat{\Pi}(t)$  of the system,  $R$  is allocated to  $J$  only if  $J$  is the job holding the resource(s) whose priority ceiling is equal to  $\hat{\Pi}(t)$ ; otherwise,  $J$ 's request is denied, and  $J$  becomes blocked.

### 3. Priority-Inheritance Rule: When $J$ becomes blocked, the job $J_i$ which blocks $J$ inherits the current priority $\pi(t)$ of $J$ . $J_i$ executes at its inherited priority until the time when it releases every resource whose priority ceiling is equal to or higher than $\pi(t)$ ; at that time, the priority of $J_i$ returns to its priority $\pi_i(t')$ at the time $t'$ when it was granted the resource(s).




# Priority-Ceiling Protocol

- PCP might deny a resource request even if the resource is free
  - PIP grants a resource request as the resource is free
- PCP imposes three types of blocking on jobs
  - Direct blocking
  - Priority-inheritance blocking
  - Avoidance blocking (new)
    - Priority-ceiling blocking
    - Actually, a kind of priority-inheritance blocking

# Priority-Ceiling Protocol


- Lemma: PCP avoids uncontrolled blocking
- How about deadlocks?



# Priority-Ceiling Protocol

- Theorem: PCP avoids deadlocks

Proof:




Jx and Jy can not  
simultaneously lock  
resources!!

# Priority-Ceiling Protocol

- Theorem: PCP avoids deadlocks

Proof:



If any one job successfully locks some resources, then at least one of the other jobs can not lock resources



# Priority-Ceiling Protocol

- Jobs governed by PCP seem having longer response times
  - Because of priority-ceiling blocking (which does not happen in PIP)
- In the worst case, PCP **does not** impose longer blocking time on jobs than PIP does
  - Recall that jobs might suffer chain blocking and transitive blocking under PIP

# Priority-Ceiling Protocol

- Theorem: any job governed by PCP can be blocked for at most the duration of one critical section

Proof:

- A job can ever be blocked by only one job
- No transitive blocking

# Priority-Ceiling Protocol


- A job can ever be blocked by only one job
  - If job J acquires resource R after blocking, then the current system ceiling is no lower than J's priority and therefore no low-priority jobs can **lock resources**
  - If job J acquires resource R, then all resources required by J have been released by other jobs

# Priority-Ceiling Protocol

- No transitive blocking
  - A job ( $J$  in the following statements) that holds some resource can not be blocked by any other jobs
    - If  $J_H$  is blocked by  $J$  when locking  $R$ , then (1)  $J$  has acquired  $R$  or (2)  $J$  has acquired  $R'$ , which's priority ceiling is higher than that of  $R$
    - Since  $J$  has acquired  $R$  or  $R'$ , then the current system priority ceiling is no lower than  $J$ 's priority. Therefore any  $J_L$  has released any resource needed by  $J$

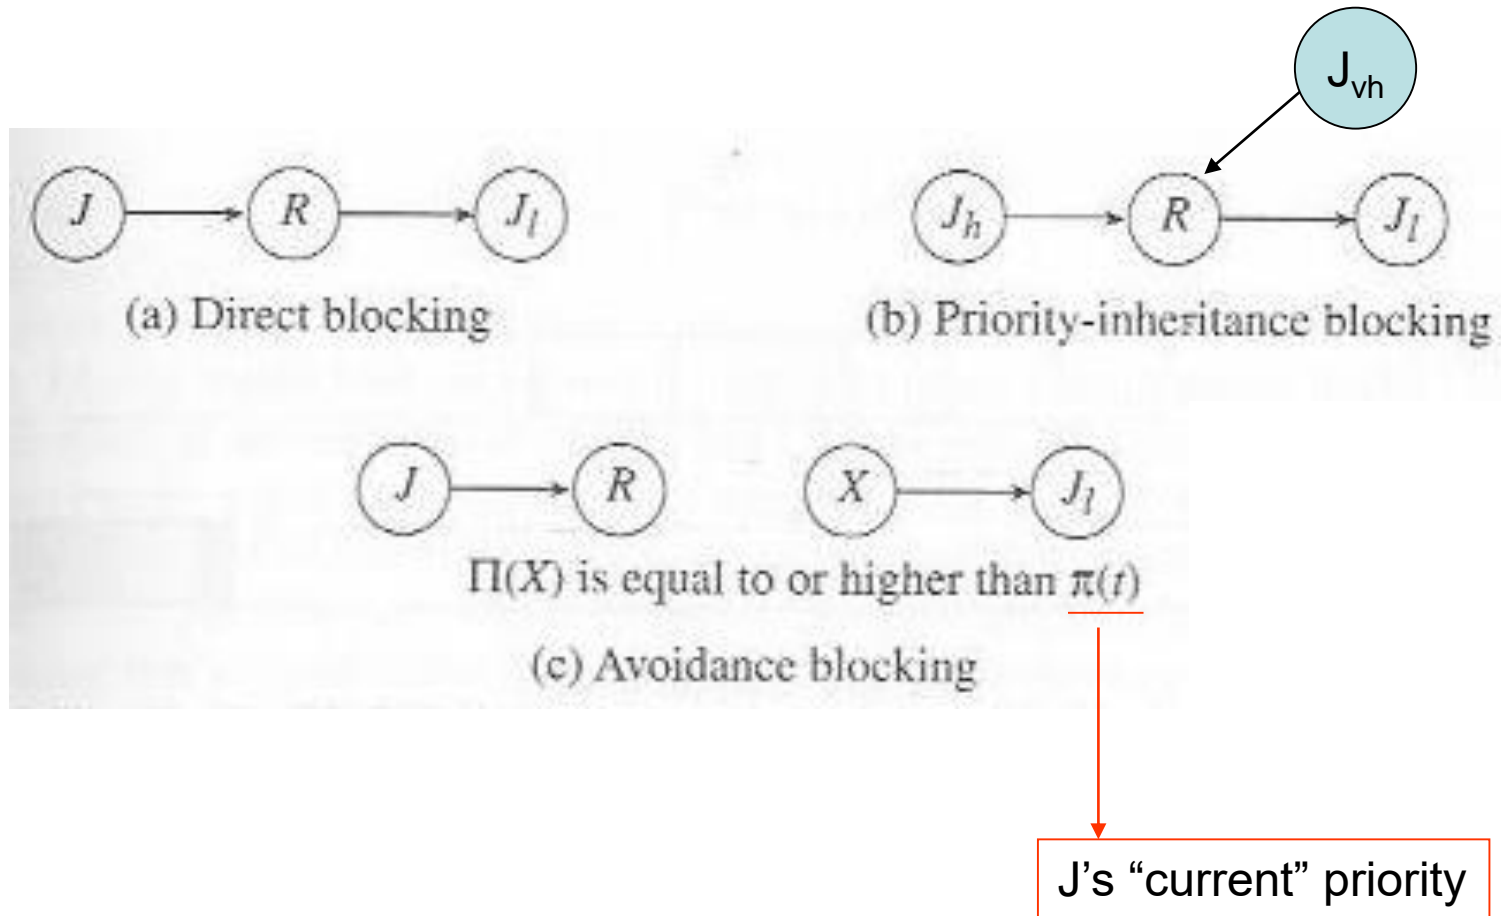
# Priority-Ceiling Protocol

- Blocking time




$J_4$  can block  $J_2$  by (1 ) priority ceiling of *Black* or (2) an priority inherited from  $J_1$

# Priority-Ceiling Protocol



# Priority-Ceiling Protocol



6,6,5,4,4

	Directly blocked by					Priority-inher blocked by					Priority-ceiling blocked by				
	$J_2$	$J_3$	$J_4$	$J_5$	$J_6$	$J_2$	$J_3$	$J_4$	$J_5$	$J_6$	$J_2$	$J_3$	$J_4$	$J_5$	$J_6$
$J_1$		6			2										
$J_2$	*		5			*	6			2	*	6			2
$J_3$		*			4		*	5		2		*	5		2
$J_4$			*					*		4			*		4
$J_5$				*					*	4				*	

# Priority-Ceiling Protocol

- For any job J that requires some resource  $\{R_1, R_2, \dots, R_n\}$ , the blocking time due to priority inheritance and that due to priority ceiling are the same
  - Priority ceiling of each resource  $R_i \geq$  any priority that J inherits from high-priority jobs
    - The highest priority that J inherits = priority ceiling of R



# Priority-Ceiling Protocol

- A collection of tasks  $\{T_1, T_2, \dots, T_n\}$  governed by PCP are schedulable if



$$\forall i, 1 \leq i \leq n, \quad \frac{b_i}{p_i} + \sum_{j=1}^i \frac{c_j}{p_j} \leq U(i)$$

or  $\forall i, 1 \leq i \leq n, \quad \max\left\{\frac{b_i}{p_i}\right\} + \sum_{j=1}^n \frac{c_j}{p_j} \leq U(n)$

- Where  $b_i$  is the longest blocking time imposed on  $T_i$

# Priority-Ceiling Protocol


$$\forall i, 1 \leq i \leq n, \quad \max\left\{\frac{b_i}{p_i}\right\} + \sum_{j=1}^n \frac{c_j}{p_j} \leq U(n)$$



The critical instant of  $T_i$  with or without blocking

# Priority-Ceiling Protocol

- Fig 8-16 Another scenario: T1, T2 and T3 arrive at epsilon, T4 arrives at 0 and blocks T1.



# Priority-Ceiling Protocol

- Context-switch issue
  - As mentioned earlier, two context switches are accounted to the preempting job
    - Thus every job is accounted to  $2 \cdot CS$
  - Under PCP, a job can be blocked for at most once
    - Thus additional  $2 \cdot CS$  is accounted to a job that potentially can “be blocked”

# Priority-Ceiling Protocol

- A collection of tasks  $\{T_1, T_2, \dots, T_n\}$  governed by PCP are schedulable if

$$\forall i, 1 \leq i \leq n, \quad \frac{b_i}{p_i} + \sum_{j=1}^i \frac{f(j) * CS + c_j}{p_j} \leq U(i)$$

in which  $f(j)=4$  if  $T_j$  's blocking time is not zero, otherwise  $f(j) = 2$ .

---

# Ceiling priority protocol

- If R is in use, T is blocked.
- If R is free, R is allocated to T. T's execution priority is raised to the priority ceiling of R if that is higher. At any given time, T's execution priority equals the highest priority ceiling of all its held resources.

# Ceiling priority protocol

- T's priority is assigned the next-highest priority ceiling of another resource when the resource with the highest priority ceiling is released.
- The task returns to its assigned priority after it has released all resources.

# Stack-Resource Policy

- Motivation
  - PCP is originally designed for RMS, but so far no resource-synchronization protocols are there for EDF
  - PCP shows high concurrency with the price that introduces extra context switches



# Stack-Resource Policy

- In EDF, no task is assigned to a fixed priority thus PCP is not directly applicable
  - EDF is a job-level fixed priority scheduling algorithm
  - Dynamic in task level
- Consider job  $J_i$  is released at time  $r_i$  and its relative deadline is  $d_i$ 
  - Job  $J_i$  could be preempted by job  $J_j$  only if  $d_j - r_j < d_i - r_i$ 
    - I.e., Job  $J_j$  could be blocked by  $J_i$

# Stack-Resource Policy

- In periodic systems, the difference between deadline and release time of jobs with respect to a task is a constant
  - In other words, jobs of a task can only be blocked by jobs of another task which has a longer period
- Because the notation “priority” does not apply to tasks in EDF, tasks are associated with “preemption level” instead
  - The **shorter** task periods are, the higher their preemption levels are

# Stack-Resource Policy

- Because jobs under PCP may be blocked after they start executing, job execution do not comply with LIFO discipline
- Job execution complies with LIFO discipline only if jobs won't voluntarily give up CPU on the way they execute
  - When a job starts executing, it executes continuously until it finishes
- If LIFO discipline is complied with, one single run-time stack is enough for tasks
  - Saves a lot of memory usage

uC/OS-II: RAM requirement = application requirement + kernel requirement + SUM(task stacks + MAX(ISR nesting))

# Stack-Resource Policy

- Basic idea of SRP
  - Task preempting / blocking in EDF is nothing different from that in RMS
    - The definition and controlling of priority inversion should be revised
  - Once a job starts executing, all resource it requires have been released by other tasks
    - To comply with LIFO discipline and to prevent deadlocks


# Stack-Resource Policy

## *Rules of Basic Stack-Based, Preemption-Ceiling Protocol*

0. *Update of the Current Ceiling:* Whenever all the resources are free, the preemption ceiling of the system is  $\Omega$ . The preemption ceiling  $\hat{\Psi}(t)$  is updated each time a resource is allocated or freed.
1. *Scheduling Rule:* After a job is released, it is blocked from starting execution until its preemption level is higher than the current ceiling  $\Psi(t)$  of the system and the preemption level of the executing job. At any time  $t$ , jobs that are not blocked are scheduled on the processor in a priority-driven, preemptive manner according to their assigned priorities.
- ★ 2. *Allocation Rule:* Whenever a job  $J$  requests for a resource  $R$ , it is allocated the resource.
3. *Priority-Inheritance Rule:* When some job is blocked from starting, the blocking job inherits the highest priority of all the blocked jobs.

**Here “priority” stands for “urgency” or “absolute deadline”!!**  
**And preemption levels reflect task periods!!**

- PCP: check system ceiling when jobs lock resources
- SRP: check system ceiling when jobs are scheduled to run



\*\* jobs of higher task preemption level do not always preempt jobs of lower preemption levels

# Stack-Resource Policy

- SRP avoids uncontrolled priority inversion
  - Protected by “priority”-inheritance
- SRP avoids transitive blocking and chain blocking
  - When a job start executing, all that resources it needs have been released, and **these resources will not be locked until the job completes**
- SRP prevent deadlocks from happening
  - A job is never be blocked once it start executing

# Stack-Resource Policy

- Blocking time
  - Direct blocking (?)
  - Deadline-inheritance blocking
  - Preemption-ceiling blocking
- Blocking time imposed on a job under SRP+EDF is the same as that of PCP+RMS



# Stack-Resource Policy

- A collection of tasks  $\{T_1, T_2, \dots, T_n\}$  governed by SRP are schedulable by EDF if

$$\forall i, \frac{b_i}{p_i} + \sum_{j=1}^n \frac{2 * CS + c_j}{p_j} \leq 1$$


- No 2 extra context switches as that in PCP!

# Stack-Resource Policy

- Pros:
  - Tasks can share one single stack
  - Easy to implement
  - Saves up 2 context switches
- Cons:
  - Relatively low concurrency, compared to PCP

# Stack-Resource Policy

- Priority ceilings for resources of multiple units
- Let  $k$  be the number of free units of resource  $R$  after some units of  $R$  is locked
- Let  $\Pi(R,k)$  be the priority ceiling (preemption ceiling) of resource  $R$  there are  $k$  units available
  - $\Pi(R,k)$  equals to the highest priority (preemption level) of all tasks that require **more than  $k$  units ( $>$ )** of resource  $R$ 
    - These task jobs may be directly blocked when trying to lock resource  $R$



Resources		$x(2)$	$y(3)$	$z(1)$
Units Required by	$J_1$	1	0	0
	$J_2$	0	2	0
	$J_3$	2	3	1
	$J_4$	1	0	1
	$J_5$	0	1	0
$\Pi(*, 0)$		$\pi_1$	$\pi_2$	$\pi_3$
$\Pi(*, 1)$		$\pi_3$	$\pi_2$	$\Omega$
$\Pi(*, 2)$		$\Omega$	$\pi_3$	$\Omega$
$\Pi(*, 3)$		$\Omega$	$\Omega$	$\Omega$

FIGURE 8-21 Example of priority ceilings of multiple-unit resources.


•Note that, the ceiling of resource R has effects to resource locking (PCP) or job scheduling (SRP) only if some units of resource R are (is) successfully locked!!

• $\Pi(R, 0)$  will be used as the ceiling of R when a resource of one unit is locked!

$R$	$N_R$	$\mu_R(1)$	$\mu_R(2)$	$\mu_R(3)$	$\lceil R \rceil_0$	$\lceil R \rceil_1$	$\lceil R \rceil_2$	$\lceil R \rceil_3$
$R_1$	3	3	2	1	3	2	1	0
$R_2$	1	1	1	0	2	0	0	0
$R_3$	3	1	3	1	3	2	2	0

Figure 4: Ceilings of Resources.

Job 1 need 3  
units of R1



Preemption levels:  $J3 > J2 > J1$

The diagram shows four horizontal bars, labeled a, b, c, and d, each representing a different data set. The x-axis is labeled from 0 to 18. Each bar is composed of colored segments (blue, green, red, yellow, white) and ends with a green circle. The segments are arranged as follows:

- Bar a: Starts at 4, ends at 13. Segments: Blue (4-6), Red (6-8), Red (8-10), Yellow (10-12), Green (12-13).
- Bar b: Starts at 2, ends at 14. Segments: Blue (2-4), Green (4-6), White (6-8), Red (8-10), Red (10-12), Green (12-14).
- Bar c: Starts at 2, ends at 16. Segments: White (2-6), Red (6-8), Red (8-10), White (10-12), White (12-14), Blue (14-16).
- Bar d: Starts at 0, ends at 17. Segments: Blue (0-2), Yellow (2-4), White (4-6), White (6-8), Yellow (8-10), Yellow (10-12), White (12-14), White (14-16), Blue (16-17).

Figure 1 shows four horizontal bars labeled a, b, c, and d, representing different configurations of a system. The x-axis is labeled from 0 to 18. Each bar is composed of colored segments: blue, red, yellow, green, and white. An upward arrow is at the start of each bar.

- Bar a: Starts at x=4, ends at x=11. Segments: blue (4-5), blue (5-6), red (6-7), red (7-8), yellow (8-9), green (9-10), blue (10-11).
- Bar b: Starts at x=2, ends at x=14. Segments: blue (2-3), red (3-4), white (4-5), white (5-6), red (6-7), red (7-8), white (8-9), white (9-10), green (10-11), green (11-12), blue (12-13), green (13-14).
- Bar c: Starts at x=2, ends at x=16. Segments: white (2-3), red (3-4), white (4-5), white (5-6), red (6-7), red (7-8), white (8-9), white (9-10), white (10-11), white (11-12), white (12-13), blue (13-14), blue (14-15), green (15-16).
- Bar d: Starts at x=0, ends at x=17. Segments: blue (0-1), yellow (1-2), white (2-3), yellow (3-4), white (4-5), white (5-6), yellow (6-7), yellow (7-8), white (8-9), white (9-10), white (10-11), white (11-12), white (12-13), white (13-14), white (14-15), blue (15-16), green (16-17).

Figure 1 shows four horizontal bars (a, b, c, d) representing different configurations of a system. The x-axis is labeled from 0 to 18. Bar a starts at x=4 and ends at x=10. Bar b starts at x=2 and ends at x=14. Bar c starts at x=2 and ends at x=16. Bar d starts at x=0 and ends at x=17. Each bar is composed of colored segments: red, blue, yellow, green, and white. Arrows indicate the direction of movement for each bar.



	Anomalies	Uncontrollable priority inversion	Bounded blocking time	Blocked at most once	Deadlock avoidance
Non-preemptible Critical Sections ( <b>NPCS</b> )	Yes	No	Yes	Yes	Yes
Ceiling-Priority Protocol	Yes	No	Yes	Yes	Yes
Priority-Inheritance Protocol ( <b>PIP</b> )	Yes	No	Yes	No	No
Priority-Ceiling Protocol ( <b>PCP</b> )	Yes	No	Yes	Yes	Yes
Stack-Resource Policy ( <b>SRP</b> )	Yes	No	Yes	Yes	Yes

# Summary

- Check list
  - Priority inversion
  - Blocking time
  - Blocking duration
  - Deadlock
  - Concurrency
  - Multiple-unit resources
- The tradeoff between concurrency and simplicity is always the primary challenge